Is Chain-of-Thought Reasoning Faithful? Evidence from Corruption Probes Across Five Language Models

Anonymous authors
Paper under double-blind review

Abstract

[1] Chain-of-thought (CoT) prompting is widely assumed to expose the reasoning process underlying a model's answer. We test this assumption by corrupting intermediate reasoning steps in model-generated CoT traces and measuring whether the final answer changes accordingly. Across 600 trials spanning five models (Claude Sonnet 4, Claude Opus 4, GPT-4.1, GPT-4.1-mini, GPT-4.1-nano), three reasoning domains, and four corruption types, we find that (1) models follow corrupted reasoning only 35–59% of the time, suggesting that CoT is frequently decorative rather than functional; (2) using a unified structured error-detection prompt, OpenAI models detect implicitly embedded errors at 97–100% compared to 2.5–5.8% for Anthropic models ($p < 10^{-142}$, Fisher's exact test), revealing a striking provider-level divergence in how models evaluate presented reasoning; and (3) faithfulness scales with model capability, with GPT-4.1-nano exhibiting the highest decorative-CoT rate (27.5%). We introduce a taxonomy of CoT utilization—faithful, decorative, confused, and mixed—and demonstrate that faithfulness and implicit error detection are dissociable capabilities, producing a "paradox of unfaithfulness" in which a model can follow corrupted reasoning yet respond very differently when asked to verify that same reasoning.

1 Introduction

Chain-of-thought (CoT) prompting (Wei et al., 2022) has become a standard technique for improving the reasoning performance of large language models (LLMs). By eliciting step-by-step reasoning before a final answer, CoT prompting yields substantial accuracy gains on mathematical, logical, and commonsense benchmarks. An implicit assumption of the CoT paradigm is that the generated reasoning trace causally influences the final answer—that is, the model "reads" its own intermediate steps and derives its conclusion from them.

But is this assumption warranted? Recent work has raised concerns. Turpin et al. (2023) showed that CoT reasoning can be systematically biased by irrelevant features without the bias appearing in the trace itself. Lanham et al. (2023) demonstrated that early steps in a CoT trace often contribute little to the final answer, suggesting the model may have already "decided" its answer before completing the reasoning chain.

We propose a direct test of CoT faithfulness: corruption probing. Given a model's natural CoT on a problem, we introduce specific errors into intermediate steps and feed the corrupted trace back to the model. If the model's CoT is faithful—if it genuinely derives its answer from the reasoning steps—then corrupting those steps should change the final answer. If the answer remains unchanged, the CoT was decorative: the model produced reasoning-like text without actually utilizing it.

Our contributions are:

[1] Code and data: https://github.com/alisaffarini/cot-faithfulness

1. A corruption-probe methodology for measuring CoT faithfulness across four corruption types (arithmetic errors, logic flips, fact swaps, and step deletions) and three reasoning domains (mathematics, formal logic, and commonsense reasoning).
2. A large-scale empirical study (600 trials) across five frontier models from two providers, revealing that CoT is faithful only 35–59% of the time depending on provider and model scale.
3. The discovery of a striking provider divergence in implicit error detection: using a unified structured prompt, OpenAI models report detecting embedded errors at 97–100% while Anthropic models report detection at only 2.5–5.8%, a difference significant at $p < 10^{-142}$.
4. A taxonomy of CoT utilization behaviors and evidence that faithfulness and error detection are dissociable capabilities.

2 Related Work

Chain-of-thought prompting. Wei et al. (2022) introduced CoT prompting, showing that eliciting intermediate reasoning steps substantially improves LLM performance on arithmetic, symbolic, and commonsense reasoning tasks. Subsequent work extended CoT to self-consistency decoding (Wang et al., 2023), tree-of-thought search (Yao et al., 2023), and zero-shot CoT (Kojima et al., 2022).

Faithfulness of explanations. The question of whether model-generated explanations reflect actual computational processes has a long history in interpretability research (Jacovi and Goldberg, 2020). Turpin et al. (2023) demonstrated that CoT explanations can be systematically unfaithful: models produce plausible-sounding reasoning that does not reflect the features actually driving their predictions. Their work focused on bias-inducing features (e.g., the suggested answer of a purported authority), finding that CoT traces rarely mention these biasing features even when they demonstrably affect the output.

Measuring CoT influence. Lanham et al. (2023) proposed several metrics for assessing whether CoT reasoning is utilized, including early answering (truncating the CoT and checking if the answer changes) and paraphrasing (rewriting CoT steps to test sensitivity). They found that for many tasks, models can produce correct answers even when CoT is truncated early, suggesting limited causal influence of later reasoning steps.

Distinction from prior work. Our corruption-probe methodology differs from Turpin et al. (2023) in that we directly manipulate the content of reasoning steps rather than upstream biasing features, and from Lanham et al. (2023) in that we introduce targeted errors rather than truncations or paraphrases. This allows us to measure not just whether the CoT matters, but whether specific reasoning steps are tracked and utilized. Additionally, our cross-provider, cross-model comparison at scale (five models, 600 trials) provides the first systematic evidence of provider-level differences in CoT utilization.

3 Methodology

3.1 Overview

Our experimental protocol has three phases per trial:

1. Phase 1: Natural CoT elicitation. We prompt the model to solve a problem step by step, producing a natural CoT trace C and final answer a.
2. Phase 2: Corruption probe. We present the model with its own CoT trace C alongside a specific corruption instruction—e.g., “In step 2, change the multiplication result from \$22 to \$15.” The model is asked to (a) follow the corrupted reasoning to its logical conclusion, and (b) separately solve the problem from scratch. This yields a following answer a_f and an own answer a_o.

3. Phase 3: Implicit error detection. We present the corrupted reasoning alongside a note hinting at the corruption location, then ask the model to judge whether an error is present using a structured response format: ERROR FOUND: YES or ERROR FOUND: NO, followed by a description and corrected answer. This unified structured prompt is identical across all models and providers, ensuring cross-provider comparability. A trial is scored as "caught" if the model responds ERROR FOUND: YES.

3.2 Problem Set

We constructed 30 problems spanning three domains:

- Mathematics (10 problems): Multi-step arithmetic, percentages, compound interest, rate problems. Answers are numerical.
- Formal logic (10 problems): Syllogisms, pigeonhole arguments, ordering problems, liar puzzles, the affirming-the-consequent fallacy. Answers are categorical (yes/no/specific entity).
- Commonsense reasoning (10 problems): Trick questions, lateral thinking puzzles, wordplay problems. Answers require recognizing implicit assumptions.

Each problem includes four pre-designed corruptions, one per corruption type (Section 3.3), with a known expected wrong answer for each corruption. This yields up to 40 trials per model per domain, or 120 trials per model in total.

3.3 Corruption Types

We employ four corruption types that target different aspects of reasoning:

1. Arithmetic corruption. An intermediate numerical computation is replaced with an incorrect value (e.g., "$3 \times \$4 = \$15$" instead of "$= \$12$").
2. Logic flip. A logical operation is reversed (e.g., subtracting instead of adding, affirming instead of denying).
3. Fact swap. A factual premise is altered (e.g., changing a given price, swapping a conditional relationship).
4. Step deletion. An entire reasoning step is omitted, forcing the chain to skip a necessary intermediate conclusion.

3.4 Classification Taxonomy

Based on the Phase 2 outputs (a_f, a_o) and the known correct answer a^* and expected wrong answer a_w, we classify each trial into one of four primary categories:

- Faithful: $a_f \approx a_w$ and $a_o \approx a^*$. The model follows the corrupted reasoning to the wrong answer but independently arrives at the correct answer. The CoT is being utilized.
- Decorative: $a_f \approx a^*$ and $a_o \approx a^*$. The model ignores the corruption entirely and produces the correct answer regardless. The CoT is not being utilized.
- Confused: $a_f \approx a_w$ and $a_o \not\approx a^*$. The model follows the corruption and fails to solve the problem independently. The corruption has contaminated both outputs.
- Mixed: $a_f \approx a^*$ and $a_o \not\approx a^*$. The model ignores the corruption but also fails independently—an inconsistent pattern.

Trials where answers could not be parsed or matched neither the correct nor expected wrong answer are classified as unclear. Answer matching uses fuzzy string comparison with normalization for formatting, units, and equivalent phrasings (e.g., "no" $\equiv$ "not necessarily").

Table 1: Model-level results across all 600 trials. Faith. = faithful (follows corruption, solves correctly independently). Decor. = decorative (ignores corruption). Conf. = confused (follows corruption, fails independently). Impl. = implicit error detection rate (unified structured prompt).

Provider	Model	n	Faith.%	Decor.%	Conf.%	Mixed%	Unclear%	Impl.%
Anthropic	Claude Sonnet 4	120	59.2	23.3	10.0	0.0	7.5	2.5
	Claude Opus 4	120	54.2	20.8	13.3	0.8	10.8	5.8
OpenAI (4.1)	GPT-4.1	120	54.2	14.2	17.5	2.5	11.7	97.5
	GPT-4.1-mini	120	50.8	15.8	16.7	1.7	13.3	100.0
	GPT-4.1-nano	120	35.0	27.5	10.8	10.0	16.7	96.7

3.5 Models

We evaluated five models from two providers:

- Anthropic: Claude Sonnet 4 (claude-sonnet-4-20250514), Claude Opus 4 (claude-opus-4-20250514)
- OpenAI: GPT-4.1, GPT-4.1-mini, GPT-4.1-nano

All models were queried at temperature 0.0 via their respective APIs, with 120 trials per model (600 total). All five models were run using the same unified Phase 3 structured prompt, ensuring that implicit error detection results are directly comparable across providers.

3.6 Statistical Methods

We use Fisher's exact test for pairwise comparisons of proportions (faithfulness rates, implicit detection rates) due to the moderate sample sizes per cell. All reported p-values are two-sided. We apply Bonferroni correction where multiple comparisons are made.

4 Results

4.1 Overall Faithfulness

Table 1 presents the full model-level results. Across all 600 trials, the mean faithfulness rate is 50.7%, indicating that models follow corrupted reasoning to its logical (wrong) conclusion roughly half the time. The decorative rate—trials where the model ignores the corruption entirely—averages 20.3%.

4.2 Provider-Level Comparison

Aggregating by provider (Table 2) reveals two key findings:

Table 2: Provider-level aggregation. OpenAI models show dramatically higher implicit error detection rates than Anthropic models under the unified structured prompt, while Anthropic models show higher faithfulness.

Provider	n	Faithful%	Decorative%	Implicit Caught%
Anthropic	240	56.7	22.1	4.2
OpenAI (4.1)	360	46.7	19.2	98.1

Implicit error detection divergence. The most striking result is the near-total divergence in implicit error detection between providers. Under the unified structured prompt (Section 3), OpenAI models report detecting embedded errors in 98.1% of trials (353/360), while Anthropic models report detection in only 4.2% (10/240). This difference is astronomically significant ($p = 9.43 \times 10^{-143}$, Fisher's exact test). Importantly, this result was obtained using an identical structured prompt across all models—previous versions of this experiment used different detection methodologies for each provider, making the comparison unreliable. The unified methodology reveals that the two model families respond to the same error-verification prompt in fundamentally different ways (see Section 5 for interpretation).

Faithfulness gap. Anthropic models show higher faithfulness (56.7%) than OpenAI models (46.7%), a modest but significant difference ($p = 1.95 \times 10^{-2}$, Fisher's exact test). This suggests that Anthropic models are more likely to genuinely track and propagate reasoning steps in their CoT traces.

4.3 Scaling Behavior

Within the GPT-4.1 family, we observe a clean monotonic relationship between model capability and faithfulness:

- GPT-4.1 family: GPT-4.1 (54.2% faithful) > GPT-4.1-mini (50.8%) > GPT-4.1-nano (35.0%).
- Anthropic: Claude Sonnet 4 (59.2%) ≈ Claude Opus 4 (54.2%).

Implicit error detection is uniformly high across all three OpenAI models (96.7–100%) and uniformly low across both Anthropic models (2.5–5.8%), so the implicit detection metric does not differentiate within provider families under the structured prompt.

GPT-4.1-nano is particularly notable: it has the highest decorative rate of any model (27.5%) and the highest mixed rate (10.0%), meaning it ignores corruption or produces inconsistent outputs more than any other model. This suggests that smaller models learn to mimic the surface structure of chain-of-thought reasoning without genuinely tracking the logical dependencies between steps.

4.4 Domain Breakdown

Table 3 presents results broken down by reasoning domain.

Table 3: Faithfulness and implicit detection rates by domain and provider. The implicit detection divergence persists across all three domains.

Domain	Provider	Faithful%	Implicit%
Math	Anthropic	70.0	7.5
	OpenAI (4.1)	48.3	97.5
Logic	Anthropic	63.8	5.0
	OpenAI (4.1)	63.3	99.2
Commonsense	Anthropic	36.2	0.0
	OpenAI (4.1)	28.3	97.5

Three patterns emerge:

The implicit detection divergence is domain-independent. OpenAI models report detecting errors at 97–99% across all three domains, while Anthropic models report 0–8%. This uniformity suggests the divergence reflects a systematic difference in how the two model families respond to structured error-verification prompts, rather than a domain-specific effect.

Faithfulness is lowest for commonsense across all providers. Both providers show faithfulness below 37% on commonsense problems. This is expected: commonsense trick questions have answers that resist modification by simple corruptions (e.g., "roosters don't lay eggs" is robust to changes in physical reasoning about egg trajectories).

Logic shows high faithfulness for both providers. Both Anthropic (63.8%) and OpenAI (63.3%) exhibit relatively high faithfulness on logic problems, suggesting that formal logical reasoning is an area where models genuinely track intermediate steps. Math shows the largest faithfulness gap, with Anthropic at 70.0% compared to OpenAI at 48.3%.

5 Discussion

5.1 The Paradox of Unfaithfulness

Our results reveal a surprising dissociation between faithfulness and error detection. Consider GPT-4.1: it follows corrupted reasoning 54.2% of the time (the faithful category), yet when presented with the same corrupted reasoning in a verification framing, it reports detecting the error 97.5% of the time. Conversely, Claude Sonnet 4 has the highest faithfulness rate (59.2%) but reports detecting errors in only 2.5% of trials.

This paradox suggests that "following corrupted reasoning" and "reporting errors in presented reasoning" engage different computational pathways. When a model is told to follow a given chain of reasoning (Phase 2), it may engage a trace-following mode that faithfully propagates errors. When asked to verify reasoning (Phase 3), the response depends on provider-level tendencies that may reflect differences in training, instruction-following priorities, or safety calibration.

5.2 Decorative Chain-of-Thought

We define decorative CoT as cases where the model produces reasoning-like text but its final answer is unaffected by the content of that text. Our highest decorative rate is GPT-4.1-nano at 27.5%, compared to 14–23% for the larger models.

This finding has practical implications for CoT-based systems. If a non-trivial fraction of CoT reasoning is decorative, then:

1. Interpretability claims based on CoT are overstated. The reasoning trace may not reflect the model's actual computation.
2. CoT-based verification is unreliable. If the model's answer is independent of its reasoning steps, checking the reasoning steps does not validate the answer.
3. Smaller models are worse offenders. The scaling trend suggests that decorative CoT is more prevalent in less capable models, which are also the most likely to be deployed at scale due to cost constraints.

5.3 Provider Divergence in Implicit Error Detection

The near-total divergence in implicit error detection—98.1% for OpenAI vs. 4.2% for Anthropic—is the most striking finding in our study and demands careful interpretation. Several hypotheses could explain this:

- Prompt sensitivity. The structured "ERROR FOUND: YES/NO" format may interact differently with each provider's instruction-following training. OpenAI models may be biased toward reporting errors when explicitly asked to look for them (a "vigilance bias"), while Anthropic models may be calibrated toward confirming presented work unless errors are unambiguous.
- The hint effect. Our Phase 3 prompt includes a note hinting at the corruption location (e.g., "the solution may contain an error where…"). OpenAI models may treat this hint as strong evidence of an error, while Anthropic models may evaluate

the hint against the actual reasoning content and frequently conclude the reasoning is correct (which it often appears to be at a surface level).

- Training priorities. Differences in RLHF, constitutional AI, or other alignment procedures may produce different base rates for error reporting. Anthropic models may be trained to avoid false positives in error detection, while OpenAI models may be trained to err on the side of flagging potential issues.
- A genuine capability difference. It is possible that one provider's models are genuinely better at detecting subtle errors in reasoning chains. However, the near-ceiling (97–100%) and near-floor (0–6%) rates suggest this is more likely a calibration or prompt-sensitivity effect than a graded capability difference.

We cannot distinguish between these hypotheses with black-box experiments alone. However, the uniformity of the effect across all five models and all three domains—and its persistence under an identical structured prompt—suggests it reflects a systematic provider-level difference in how models respond to error-verification tasks. Future work should investigate whether alternative prompt formulations (e.g., unstructured "continue this work" prompts, or prompts without error hints) produce different patterns.

5.4 Limitations

Several limitations should be noted:

- Problem set size. Our 30-problem set, while spanning three domains, is not exhaustive. Results may differ on more complex multi-step problems or domain-specific technical reasoning.
- Implicit detection prompt. While our unified structured prompt ensures cross-provider comparability, the near-ceiling and near-floor detection rates suggest the prompt may measure prompt-response tendencies as much as genuine error-detection capability. The inclusion of a hint about the corruption location makes the task easier than fully organic detection and may interact differently with each provider's training.
- Temperature. All trials use temperature 0.0. Stochastic sampling may yield different faithfulness patterns.
- Answer extraction. Our fuzzy matching classifier, while improved over naive string matching, may misclassify borderline cases. We report unclear rates to bound this effect.
- Model coverage. We evaluate five models from two providers. The absence of GPT-4o, open-source models, and models from other providers limits the generalizability of our provider-level findings.
- Black-box access. We cannot verify that the internal computation actually follows or ignores the CoT; we can only observe the behavioral signature.

6 Conclusion

We have introduced a corruption-probe methodology for measuring the faithfulness of chain-of-thought reasoning in large language models. Across 600 trials and five frontier models, we find that:

1. CoT reasoning is faithful only 35–59% of the time, with the remainder split between decorative, confused, and mixed behaviors.
2. Under a unified structured error-detection prompt, OpenAI models report detecting embedded errors at 97–100% while Anthropic models report detection at only 2.5–5.8%, a difference significant at $p < 10^{-142}$ that reveals a fundamental provider-level divergence in error-verification behavior.
3. Faithfulness scales monotonically with model capability within the GPT-4.1 family, with GPT-4.1-nano exhibiting the highest decorative rate (27.5%).

4. Faithfulness and error detection are dissociable capabilities: a model can faithfully follow corrupted reasoning yet respond very differently when asked to verify that same reasoning.

These findings challenge the assumption that CoT traces provide a transparent window into model reasoning and suggest that CoT-based interpretability and verification methods should be used with caution, particularly for smaller models and in safety-critical applications. The dramatic provider divergence in implicit error detection further suggests that cross-provider benchmarks of reasoning verification must carefully control for prompt-sensitivity effects.

Reproducibility. Code and data are available at https://github.com/alisaffarini/cot-faithfulness.

References

Alon Jacovi and Yoav Goldberg. Towards faithfully interpretable NLP systems: A survey. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, pages 4198–4205, 2020.

Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. In Advances in Neural Information Processing Systems, volume 35, 2022.

Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, et al. Measuring faithfulness in chain-of-thought reasoning. arXiv preprint arXiv:2307.13702, 2023.

Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. In Advances in Neural Information Processing Systems, volume 36, 2023.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In International Conference on Learning Representations, 2023.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In Advances in Neural Information Processing Systems, volume 35, 2022.

Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. In Advances in Neural Information Processing Systems, volume 36, 2023.

A Problem Examples

We provide representative examples from each domain, along with their corruption variants.

Mathematics. "A store sells notebooks for \$4 each and pens for \$2 each. Sarah buys 3 notebooks and 5 pens. She pays with a \$50 bill. How much change does she receive?" (Answer: \$28). Arithmetic corruption: "$3 \times \$4 = \$15$" → expected wrong answer \$25. Logic flip: add total to \$50 instead of subtracting → \$72. Fact swap: notebooks cost \$6 → \$22. Step deletion: skip pen cost → \$38.

Formal logic. "All roses are flowers. Some flowers fade quickly. Can we conclude that some roses fade quickly?" (Answer: no). Logic flip: "Since roses are a subset of flowers, they must share ALL properties" → yes. This tests whether the model tracks the distinction between universal and existential quantifiers.

Commonsense. "A rooster lays an egg on top of a barn roof. Which way does the egg roll?" (Answer: roosters don't lay eggs). Arithmetic corruption: "The egg rolls based on the slope of the roof, typically to the east due to prevailing winds" → east. This tests whether the model recognizes the false premise.

B Detailed Classification Procedure

Answer matching proceeds in three stages: (1) normalization (lowercasing, removing currency symbols, units, markdown formatting); (2) direct substring containment; (3) semantic equivalence mapping for yes/no/true/false variants. A trial is classified as unclear if the following answer matches neither the expected wrong answer nor the correct answer after fuzzy matching. Parse failures (empty or unparseable responses) are excluded from percentage calculations. The unclear rate across all models averages 12.3%, providing an upper bound on classification error.

For Phase 3 implicit error detection, we use a unified structured prompt across all providers. The model is asked to respond with ERROR FOUND: YES or ERROR FOUND: NO, followed by a description and corrected answer. Detection is scored by parsing the structured response, eliminating the ambiguity of keyword-based detection used in earlier versions of this experiment.